

Lecture Three: How to Learn





Yi Ma

Director of the School of Computing and Data Science

Director of the Institute of Data Science

What to Learn? Predictable information from data sensed of the external world

Mathematically, all predictable information can be modeled as certain **low-dimensional structures** in the high-dimensional data



Computational complexity associated with realizing intelligence:
incomputable -> computable -> tractable -> scalable

Kolmogorov
& Solomonoff



Turing &
Shannon

NP vs P

DNN and BP

How to Learn? Pursuing parsimony

Data lie on a low-dim linear **subspace**



If we view the data (image) as a matrix

$$\mathbf{A} = [\mathbf{a}_1 \mid \cdots \mid \mathbf{a}_n] \in \mathbb{R}^{m \times n}$$

then

$$r \doteq \text{rank}(\mathbf{A}) \ll m.$$

Principal Component Analysis (PCA) via singular value decomposition (SVD):

- Optimal estimate of \mathbf{A} under iid Gaussian noise $\mathbf{D} = \mathbf{A} + \mathbf{Z}$
- Efficient and scalable computation
- Fundamental statistical tool, with huge impact in practice...

How to Learn? Pursuing parsimony

Singular Value Decomposition (SVD)

Given $A \in \mathbb{R}^{m \times n}$ with $\text{rank}(A) = r$, we like to decompose it into a special **matrix** form:

$$U_r = [\vec{u}_1, \vec{u}_2, \dots, \vec{u}_r] \text{ orthogonal}$$

$$V_r = [\vec{v}_1, \vec{v}_2, \dots, \vec{v}_r] \text{ orthogonal}$$

$$\Sigma_r = \text{diag}\{\sigma_1, \sigma_2, \dots, \sigma_r\} > 0$$

$$A = U_r \Sigma_r V_r^\top = [\vec{u}_1, \vec{u}_2, \dots, \vec{u}_r] \begin{bmatrix} \sigma_1 & 0 & \cdots & 0 \\ 0 & \sigma_2 & \ddots & \vdots \\ \vdots & \ddots & \ddots & 0 \\ 0 & \cdots & 0 & \sigma_r \end{bmatrix} \begin{bmatrix} \vec{v}_1^\top \\ \vec{v}_2^\top \\ \vdots \\ \vec{v}_r^\top \end{bmatrix}$$

How to Learn? Pursuing parsimony

Singular Value Decomposition (SVD)

Given $A \in \mathbb{R}^{m \times n}$ with $\text{rank}(A) = r$, we like to decompose it into a special **matrix** form:

$$A = U_r \Sigma_r V_r^\top = [\vec{u}_1, \vec{u}_2, \dots, \vec{u}_r] \begin{bmatrix} \sigma_1 & 0 & \cdots & 0 \\ 0 & \sigma_2 & \ddots & \vdots \\ \vdots & \ddots & \ddots & 0 \\ 0 & \cdots & 0 & \sigma_r \end{bmatrix} \begin{bmatrix} \vec{v}_1^\top \\ \vec{v}_2^\top \\ \vdots \\ \vec{v}_r^\top \end{bmatrix}$$

Theorem: given $A \in \mathbb{R}^{m \times n}$ with $\text{rank}(A) = r$, let $A^\top A = \sum_{i=1}^r \lambda_i \vec{v}_i \vec{v}_i^\top$ and $\sigma_i = \sqrt{\lambda_i}$, $\vec{u}_i = \frac{1}{\sigma_i} A \vec{v}_i \in \mathbb{R}^m$, $i = 1, \dots, r$. Then we have $U_r = [\vec{u}_1, \vec{u}_2, \dots, \vec{u}_r]$ orthogonal, and

$$A = \sum_{i=1}^r \sigma_i \vec{u}_i \vec{v}_i^\top = U_r \Sigma_r V_r^\top \quad \Sigma_r = \text{diag}\{\sigma_1, \dots, \sigma_r\} = \begin{bmatrix} \sigma_1 & 0 & 0 \\ 0 & \ddots & 0 \\ 0 & 0 & \sigma_r \end{bmatrix}$$

How to Learn? Pursuing parsimony

Low-Rank Approximation: Eckart-Young Theorem

Approximate a matrix $A \in \mathbb{R}^{m \times n}$ with rank $r \leq \min\{m, n\}$ by a lower-rank matrix.

$$A = [\vec{a}_1, \vec{a}_2, \dots, \vec{a}_n] = \sum_{i=1}^r \sigma_i \vec{u}_i \vec{v}_i^\top = \sum_{i=1}^{\ell} \sigma_i \vec{u}_i \vec{v}_i^\top + \sum_{i=\ell+1}^r \sigma_i \vec{u}_i \vec{v}_i^\top \quad \text{with } \sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r \geq 0$$

Theorem [Eckart-Young 1936]: The optimal solution to the low-rank approximation problem:

$$\min_{B \in \mathbb{R}^{m \times n}} \|A - B\|_F^2 \quad \text{subject to} \quad \text{rank}(B) = \ell$$

is given by: $B_\star = A_\ell = \sum_{i=1}^{\ell} \sigma_i \vec{u}_i \vec{v}_i^\top$.

How to Learn? Pursuing parsimony

Low-Rank Approximation: Rank Minimization

Approximate a matrix $A \in \mathbb{R}^{m \times n}$ with rank $r \leq \min\{m, n\}$ by a lower-rank matrix.

$$A = [\vec{a}_1, \vec{a}_2, \dots, \vec{a}_n] = \sum_{i=1}^r \sigma_i \vec{u}_i \vec{v}_i^\top = \sum_{i=1}^{\ell} \sigma_i \vec{u}_i \vec{v}_i^\top + \sum_{i=\ell+1}^r \sigma_i \vec{u}_i \vec{v}_i^\top$$

Rank minimization problem:

$$\min_{B \in \mathbb{R}^{m \times n}} \text{rank}(B) \quad \text{subject to} \quad \|A - B\|_F^2 \leq \epsilon^2 ?$$

How to Learn? Pursuing parsimony

Low-Rank Approximation: Model Selection


Approximate a matrix $A \in \mathbb{R}^{m \times n}$ with rank $r \leq \min\{m, n\}$ by a lower-rank matrix.

$$A = [\vec{a}_1, \vec{a}_2, \dots, \vec{a}_n] = \sum_{i=1}^r \sigma_i \vec{u}_i \vec{v}_i^\top = \sum_{i=1}^{\ell} \sigma_i \vec{u}_i \vec{v}_i^\top + \sum_{i=\ell+1}^r \sigma_i \vec{u}_i \vec{v}_i^\top$$

Selecting a good tradeoff between rank and residual:

1. $\min_{B \in \mathbb{R}^{m \times n}} \text{rank}(B) = d \quad \text{subject to} \quad \sigma_{d+1}^2 \leq \tau?$

2. $\min_{B \in \mathbb{R}^{m \times n}} \alpha \cdot \text{rank}(B) + \beta \cdot \sigma_{d+1}^2?$




How to Learn? Pursuing parsimony

Principal Component Analysis (Statistics)

Problem [Pearson, 1901, Hotelling, 1933]: given

$$A = [\vec{a}_1, \vec{a}_2, \dots, \vec{a}_n] \in \mathbb{R}^{m \times n} \quad \vec{\mu} = \frac{1}{n}(\vec{a}_1 + \vec{a}_2 + \dots + \vec{a}_n) = \mathbf{0}$$


find a normal vector $\|\vec{u}\|_2 = 1$ such that $\max_{\vec{u}} \|\vec{u}^\top A\|_2^2 = \|\vec{u}\vec{u}^\top A\|_2^2$.



How to Learn? Pursuing low-dimensional models

Generalized Principal Component Analysis (GPCA):

The data are on a mixture of subspaces




How to Learn? We compress to learn

1. Clustering Mixed Data (Interpolation)

Assume data $X = [x_1, x_2, \dots, x_m]$
are i.i.d. samples from a mixture
of distributions: $p(x, \mathcal{X}) = \sum_{j=1}^k \hat{\pi}_j p_j(x, \mathcal{X})$.

Classic approaches to cluster the data:
the maximum-likelihood (ML) estimate
via Expectation Maximization (EM):



$$\max_{\hat{\pi}, \hat{\mu}} E \log \sum_{j=1}^k \hat{\pi}_j p_j(x, \mathcal{X}) \quad \Leftrightarrow \quad \max_{\hat{\pi}, \hat{\mu}} \frac{1}{m} \sum_{i=1}^m \log \sum_{j=1}^k \hat{\pi}_j p_j(x_i, \mathcal{X}) .$$

Difficulties: ML is not well-defined when distributions are degenerate.


How to Learn? We compress to learn

Clustering via Compression

[Yi Ma, Harm Derksen, Wei Hong, and John Wright, TPAMI'07]

A Fundamental Idea:
Data belong to mixed low-dim
structures should be compressible.

Cluster Criterion:
Whether the number of binary bits
required to store the data is less
(information gain):



$$\# \text{ bits}(X \cup Y) \geq \# \text{ bits}(X) + \# \text{ bits}(Y)?$$

“The whole is greater than the sum of the parts.”
– Aristotle, 320 BC

How to Learn? We compress to learn


Theorem (Ma, TPAMI'07)

The number of bits needed to encode data $X = [x_1, x_2, \dots, x_m] \in \mathbb{R}^{D \times m}$ up to a precision $\|x - \hat{x}\|_2 \leq \epsilon$ is bounded by:


$$L(X, \epsilon) \leq \frac{m + D}{2} \log \det \left(I + \frac{D}{m} X^T X \right).$$

This can be derived from constructively quantifying SVD of X or by sphere packing $\text{vol}(X)$ as samples of a noisy Gaussian source.

Linear subspace



Gaussian




How to Learn? We compress to learn


Cluster to Compress

$$L(X) \geq L^c(X) \doteq L(X_1) + L(X_2) + H(|X_1|, |X_2|)?$$

partitioning:



sifting:



How to Learn? We compress to learn

A Greedy Algorithm

Seek a partition of the data $\mathbf{X} \rightarrow [\mathbf{X}_1, \mathbf{X}_2, \dots, \mathbf{X}_k]$ such that

$$\min L^c(\mathbf{X}) \doteq L(\mathbf{X}_1) + \dots + L(\mathbf{X}_k) + H(|\mathbf{X}_1|, \dots, |\mathbf{X}_k|).$$


Optimize with a *bottom-up pair-wise* merging algorithm [Ma, TPAMI'07]:

- 1: **input:** the data $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_m] \in \mathbb{R}^{D \times m}$ and a distortion $\epsilon^2 > 0$.
- 2: initialize \mathcal{S} as a set of sets with a single datum $\{S = \{\mathbf{x}\} \mid \mathbf{x} \in \mathbf{X}\}$.
- 3: **while** $|\mathcal{S}| > 1$ **do**
- 4: choose distinct sets $S_1, S_2 \in \mathcal{S}$ such that
$$L^c(S_1 \cup S_2) - L^c(S_1, S_2) \text{ is minimal.}$$
- 5: **if** $L^c(S_1 \cup S_2) - L^c(S_1, S_2) \geq 0$ **then** break;
- 6: **else** $\mathcal{S} := (\mathcal{S} \setminus \{S_1, S_2\}) \cup \{S_1 \cup S_2\}$.
- 7: **end**
- 8: **output:** \mathcal{S}



How to Learn? We compress to learn

Surprisingly Good Performance

Empirically, **find global optimum and extremely robust to outliers**



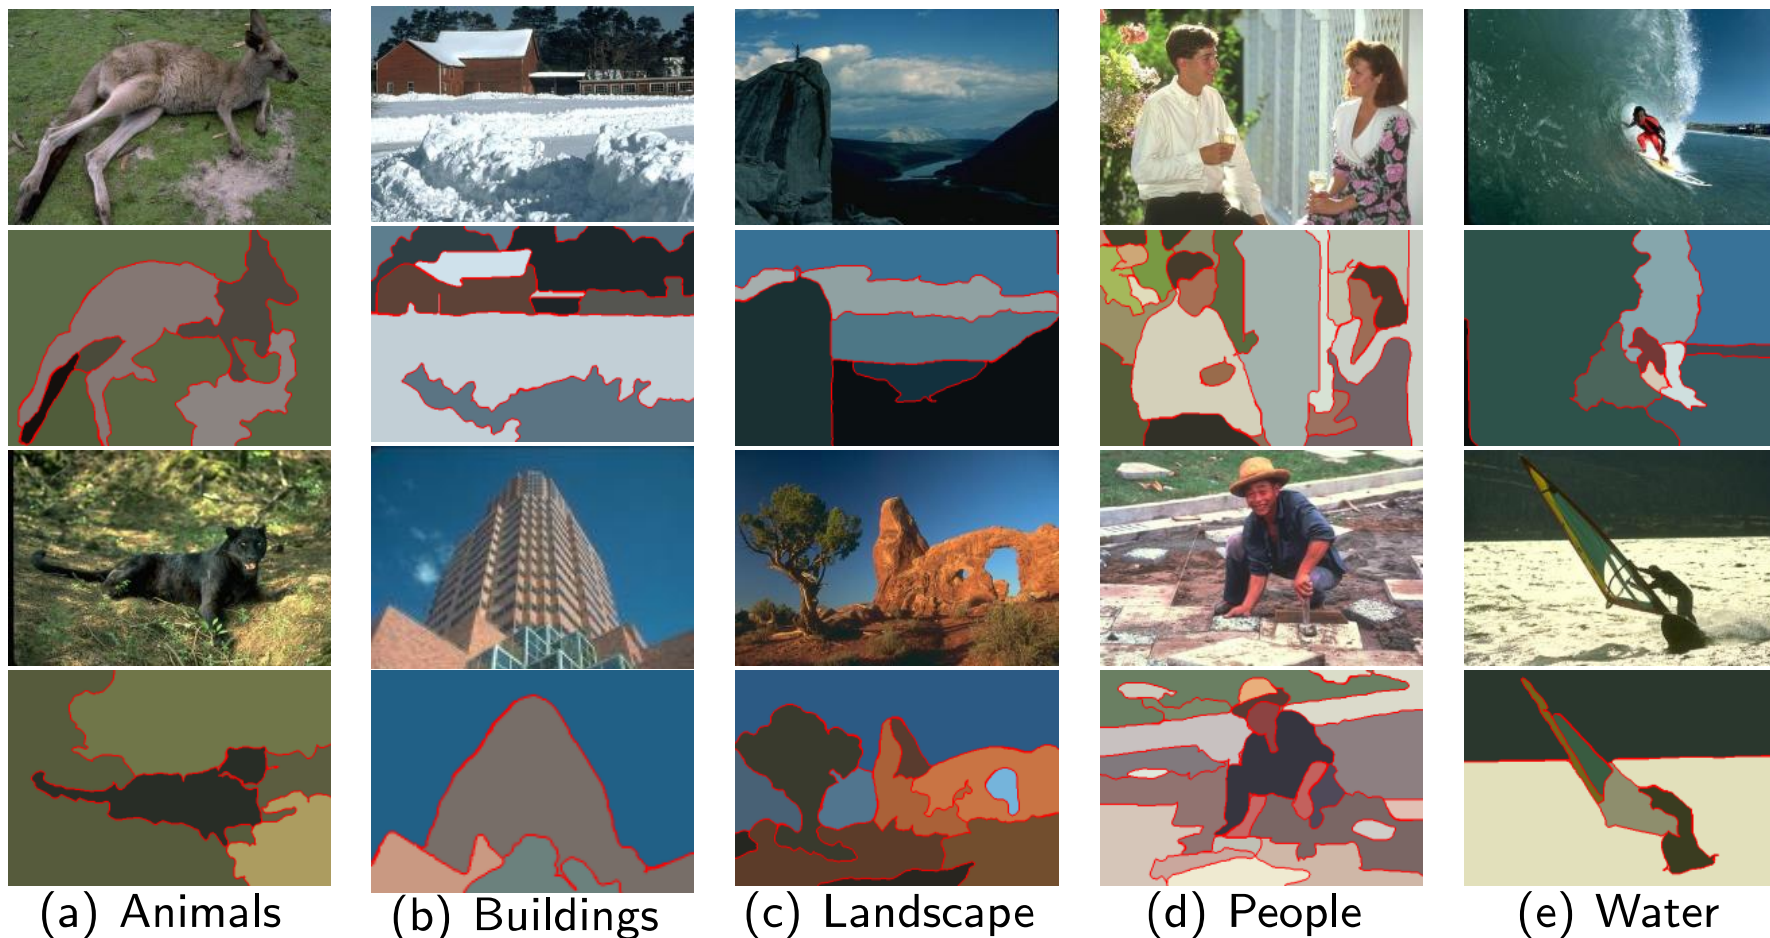
A strikingly sharp **phase transition** w.r.t. quantization ϵ





How to Learn? We compress to learn

Natural Image Segmentation [Mobahi et.al., IJCV'09]

Compression alone, without any supervision, leads to **state of the art** segmentation on natural images (and many other types of data).



How to Learn a more general low-dim distribution?



Nonlinear Manifolds